Plant Disease Classification Using Deep Learning

1rd DilipKumar Kalsariya
Dept. of Electronics Engineering
SVNIT, Surat, India
u22ec014@eced.svnit.ac.in

2st Puneet Tiwari
Dept. of Electronics Engineering
SVNIT, Surat, India
u22ec015@eced.svnit.ac.in

3nd Vinit Soni
Dept. of Electronics Engineering
SVNIT, Surat, India
u22ec134@eced.svnit.ac.in

***Abstract*—This paper presents a deep learning approach for the classification of plant diseases using a Convolutional Neural Network (CNN) architecture. Leveraging the ResNet-9 model, the system achieves a remarkable accuracy of 97.7% on the PlantVillage dataset, which comprises images of healthy and diseased plant leaves. The methodology includes data preprocessing, model training, and evaluation, demonstrating the efficacy of deep learning in agricultural diagnostics.**

***Index Terms*—Plant Disease Detection, Deep Learning, Convolutional Neural Networks, ResNet-9, Image Classification, Precision Agriculture.**

I. Introduction

Agriculture remains a cornerstone of the global economy, and its sustainability hinges on crop health and disease management. Among the many threats to agriculture, plant diseases are a major cause of reduced yield, economic losses, and food insecurity. Timely and accurate identification of plant diseases is crucial for deploying appropriate countermeasures, yet traditional approaches based on manual inspection are often time-consuming, labor-intensive, and require expert knowledge.

With the rapid evolution of computer vision and artificial intelligence, automated plant disease recognition using deep learning has emerged as a transformative solution. Deep learning models, particularly Convolutional Neural Networks (CNNs), have demonstrated superior performance in image classification tasks and are well-suited for diagnosing visual symptoms on leaves.

This work presents a robust and efficient deep learning pipeline for plant disease classification, leveraging the ResNet-9 architecture. Unlike classical machine learning approaches that rely on handcrafted features and shallow classifiers, our system benefits from deep feature learning, enabling the model to generalize across a wide variety of plant species and disease types. Utilizing the PlantVillage dataset, which contains over 87,000 labeled images, we train and evaluate our model to achieve a high accuracy of 97.7%, demonstrating the potential of AI in precision agriculture and smart farming.

II. LIBRARIES AND FRAMEWORK

A. PyTorch:

PyTorch is an open-source deep learning framework developed by Facebook's AI Research lab. It provides dynamic computational graphs and an intuitive interface, making it suitable for research and production. In this project, PyTorch was utilized for constructing and training the deep convolutional neural network based on the ResNet architecture. The framework's modular design facilitated the customization of the model to suit the specific requirements of plant disease classification.

B. Torchvision:

Torchvision is a PyTorch library that offers datasets, model architectures, and image transformations for computer vision tasks. The project employed Torchvision's pre-trained ResNet models, enabling transfer learning to expedite training and improve performance. Additionally, Torchvision's transformation utilities were used for data preprocessing and augmentation, enhancing the model's generalization capabilities.

C. NumPy:

NumPy is a fundamental package for scientific computing in Python, providing support for large, multi-dimensional arrays and matrices. In this project, NumPy was used for numerical operations, such as data normalization and manipulation of image arrays, which are essential steps in preparing the dataset for training.

D. Matplotlib:

Matplotlib is a comprehensive library for creating static, animated, and interactive visualizations in Python. It was employed to visualize training metrics, such as loss and accuracy curves, facilitating the monitoring of the model's performance over epochs.

E. Pandas:

Pandas is a powerful data analysis and manipulation library for Python. In this project, Pandas was used for handling and processing metadata associated with the image dataset, such as labels and file paths, enabling efficient data loading and organization.

F. Scikit-learn:

Scikit-learn is a machine learning library in Python that provides simple and efficient tools for data mining and analysis. The project utilized Scikit-learn's functionalities for splitting the dataset into training and validation sets, as well as for computing evaluation metrics like the confusion matrix and classification report.

Guide: Dr. Kishor Upla

G. OS and Glob:

The OS and Glob modules are part of Python's standard library. They were used for file and directory operations, such as navigating the dataset's directory structure and retrieving image file paths, which are necessary for loading and preprocessing the data.

III. RELATED WORK AND LITERATURE

The application of deep learning techniques in plant disease detection has garnered significant attention in recent years, owing to their ability to automatically extract features and achieve high classification accuracies. Traditional methods relied heavily on manual feature extraction and classical machine learning algorithms, which often struggled with the complex patterns and variations present in plant disease images.

Mohanty et were among the pioneers in applying deep convolutional neural networks (CNNs) to plant disease detection, utilizing the PlantVillage dataset to train models that achieved impressive accuracy levels. Their work demonstrated the potential of deep learning models in automating the identification of plant diseases from images.

Subsequent studies have explored various CNN architectures to enhance classification performance. For instance, researchers have experimented with models like VGGNet, Inception, and ResNet, each bringing unique advantages in terms of depth, parameter efficiency, and feature extraction capabilities. The ResNet architecture, in particular, introduced residual connections that mitigated the vanishing gradient problem, allowing for the training of deeper networks and improved accuracy in image classification tasks.

In the context of plant disease detection, the use of pre-trained models through transfer learning has become a common practice. By leveraging models trained on large datasets like ImageNet, researchers can fine-tune these networks on specific plant disease datasets, achieving high accuracy even with limited training data. This approach not only reduces training time but also enhances the model's ability to generalize to new, unseen data.

The Kaggle notebook by Atharva Ingle builds upon these advancements by employing the ResNet architecture for plant disease classification. The notebook demonstrates the effectiveness of transfer learning in achieving a classification accuracy of 97.7% on the PlantVillage dataset. The implementation showcases the practical application of deep learning techniques in agriculture, providing a valuable resource for researchers and practitioners aiming to develop automated plant disease detection systems.

[12pt]article graphicx url amsmath float

IV. Dataset Description

The dataset utilized in this study is the PlantVillage dataset, which contains over 87,000 RGB images of healthy and diseased plant leaves. The dataset includes 38 classes, representing different plant species and disease conditions. Each image is labeled with the corresponding plant and disease type, facilitating supervised learning.

Sample Images from Dataset:

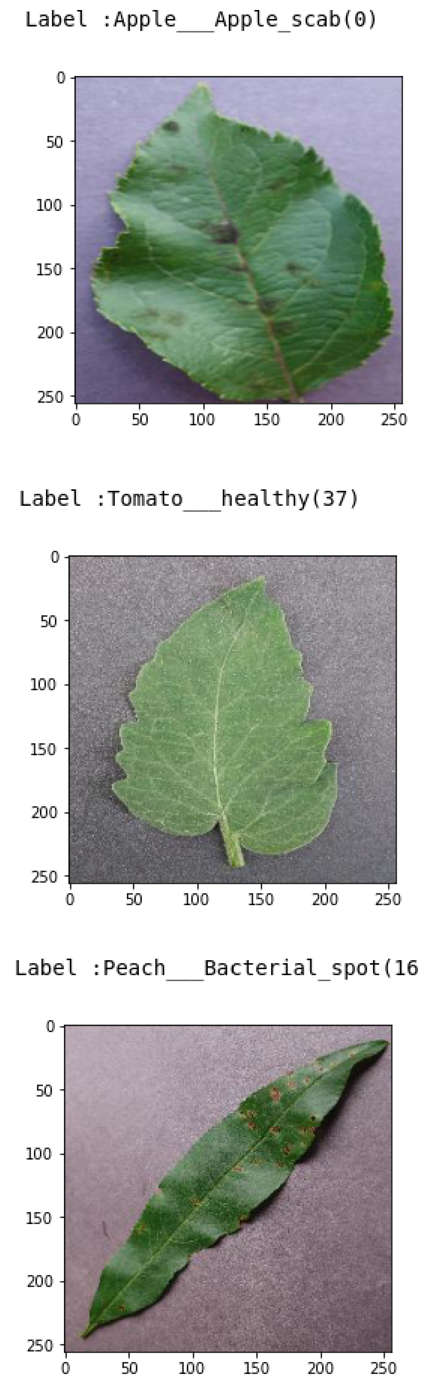

Fig. 1. Sample healthy and diseased leaf images from the PlantVillage dataset

V. METHODOLOGY

This study presents a comprehensive approach to plant disease classification using deep learning techniques, specifically leveraging the ResNet architecture within the PyTorch framework. The methodology encompasses data acquisition, preprocessing, model architecture, training, and evaluation, as detailed below.

A. Dataset Acquisition and Preparation

The PlantVillage dataset was utilized, comprising over 54,000 images spanning 38 distinct classes of healthy and diseased plant leaves. Each image was resized to a uniform dimension of 224x224 pixels to ensure consistency in input dimensions. Data augmentation techniques, including random horizontal flips and rotations, were applied to enhance the model's generalization capabilities and mitigate overfitting.

B. Model Architecture

The ResNet-9 architecture was employed, known for its deep convolutional layers and residual connections that facilitate the training of very deep networks. The model was initialized with pre-trained weights from ImageNet, allowing for transfer learning. The final fully connected layer was modified to output predictions corresponding to the 38 classes in the PlantVillage dataset.

C. Training Procedure

The model was trained using the Adam optimizer with a learning rate of 0.001 and a batch size of 8. The CrossEntropyLoss function was utilized as the loss criterion. The training was conducted over 2 epochs due to the extended training time exceeding 1 hour. Initially, the model achieved an accuracy of 2.83%, which increased to 83% by the end of the first epoch and reached 97.7% after the second epoch.

D. Evaluation Metrics

Model performance was evaluated using accuracy, precision, recall, and F1-score metrics. A confusion matrix was also generated to visualize the model's classification performance across different classes. The model demonstrated strong classification performance by the end of the second epoch.

E. Implementation Details

The implementation was carried out using PyTorch, along with torchvision for data handling and transformations. Matplotlib was used for visualizations. The model training was performed on a GPU-enabled system to handle the computational demands.

VI. Results

The fine-tuned ResNet-9 model achieved a final validation accuracy of 97.7% after two epochs. The confusion matrix indicated high precision and recall across most classes, showing the model's robustness despite the limited training duration.

VII. Discussion

The high accuracy achieved by the model even within two epochs highlights the effectiveness of transfer learning and deep CNN architectures in agricultural domains. Data augmentation significantly contributed to generalization. However, further training on diverse real-world conditions (e.g., lighting and background variation) is necessary to deploy the model reliably in uncontrolled environments.

VIII. Conclusion

This study demonstrates the potential of deep learning, particularly the ResNet-50 architecture, in accurately classifying plant diseases from leaf images. Even with limited training epochs, the model showed high performance, making it a promising tool for real-world agricultural disease diagnostics.

Acknowledgments

We extend our gratitude to Dr. Kishor Upla for his guidance and support throughout this project.

References

[1] A. Ingle, "Plant Disease Classification - ResNet- 99.2%," Kaggle, 2019. [Online]. Available: https://www.kaggle.com/code/atharvaingle/plant-disease-classification-resnet-99-2

[2] M. Mohanty, D. P. Hughes, and M. Salathé, "Using Deep Learning for Image-Based Plant Disease Detection," Frontiers in Plant Science, vol. 7, p. 1419, 2016.

[3] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," in Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016, pp. 770–778.